

INFOH415 - Advanced Databases

Jose Antonio Lorenzo Abril

Fall 2022



Professor: Gilles Dejaegere

Student e-mail: jose.lorenco.abril@ulb.be

Contents

I	Active Databases	4
1	Session 1: Exercises 'PhD'	4
2	Sessions 2 and 3: Exercises 'Employees department projects'	7
II	Graph databases	15
3	Session 4: Relational graphs	15
4	Sessions 5 and 6: Neo4j and Cypher	16
4.1	Session 5: Cypher Part I	16
4.2	Session 6: Cypher Part II	19
III	Temporal databases	27
IV	Spatial databases	28
5	Lab 10	28
6	Lab 11	30

List of Figures

List of Tables

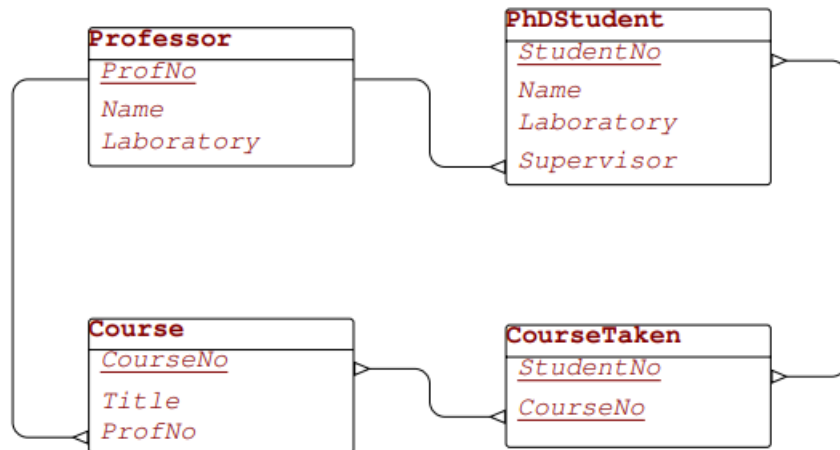
Listings

Part I

Active Databases

1 Session 1: Exercises 'PhD'

Consider the following database schema:



Define in SQL Server a set of triggers that ensure the following constraints:

1. A PhD student must work in the same laboratory as his/her supervisor.

This can be violated by:

- (a) Insert into PhDStudent.
- (b) Update of Laboratory or Supervisor in PhDStudent.

```

1  -- For a) and b)
2  -- Case Abort:
3  create trigger StudSameLabAsSuperv_PhDStud_InsUpd_Abort on PhDStudent after insert,
4  update as
5  if exists (
6      select *
7      from Inserted I, Professor P
8      where P.ProfNo = I.Superivisor
9            and P.Laboratory <> I.Laboratory )
10 begin
11     raiserror ('Constraint Violation: A PhD student must work in the same laboratory as
12             his/her supervisor', 1, 1)
13     rollback
14 end
15
16 -- Case Repair:
17 create trigger StudSameLabAsSuperv_PhDStud_InsUpd_Repair on PhDStudent after insert,
18 update as
19 begin
20     update PhDStudent
21     set Laboratory = (
22         select P.Laboratory
23         from Professor P
24         where P.ProfNo = Supervisor )
25     where StudentNo in (
26         select I.StudentNo
27         from Inserted I )
28 end
  
```

- (c) Update of Laboratory in Professor.

```

1  -- Abort
2  create trigger StudSameLabAsSuperv_Prof_Upd_Abort on Professor after update as
3  if exists (
4      select *
5      from Inserted I, PhDStudent S
6      where I.ProfNo = S.Supervisor
7            and I.Laboratory <> S.Laboratory )
8  begin
9      raiserror ('Constraint Violation: A PhD student must work in the same laboratory as
10         his/her supervisor', 1, 1)
11      rollback
12  end
13
14 -- Repair
15 create trigger StudSameLabAsSuperv_Prof_Upd_Repair on Professor after update as
16 begin
17     update PhDStudent
18     set Laboratory = (
19         select I.Laboratory
20         from Inserted I
21         where Supervisor = I.ProfNo )
22     where Supervisor in (
23         select I2.ProfNo
24         from Inserted I2 )
25 end

```

- (d) Delete from Professor: The professor is deleted and the attributes Laboratory and Supervisor of the PhD students who worked for the deleted professor are set to null.

```

1  alter table PhDStudent
2  drop constraint FK_PhDStudent_Professor -- Because SQL server r does not implement
3      the option on delete set null for the referential integrity, it is necessary to
4      drop the foreign key constraint in the table PhDStudent.
5
6  create trigger StudSameLabAsSuperv_Prof_Del_Repair on Professor after delete as
7  begin
8      update PhDStudent
9      set Laboratory = null, Supervisor = null
10     where Supervisor in (
11         select ProfNo
12         from Deleted )
13 end

```

2. A PhD student must take at least one course.

This can be violated by:

- (a) Insert into PhDStudent

```

1  -- This does not work in SQL Server, since a trigger is executed immediately after
2  -- the triggering instruction. Thus, embedding several inserts (into PhDStudent and
3  -- CourseTaken) into one transaction would not help. Practically, thus, and without
4  -- any further assumption, it will not be possible to ensure that this constraint
5  -- verified, as the aborting trigger would prevent any insertion into the table
6  -- PhDStudent.
7
8  create trigger PhDStudMinOneCourse_PhDStud_Ins_Abort on PhDStudent after insert as
9  if exists (
10     select *
11     from Inserted I
12     where not exists (
13         select *
14         from CourseTaken
15         where StudentNo = I.StudentNo ) )
16 begin
17     raiserror ('Constraint Violation: A PhD student must take at least one course', 1,
18         1)
19     rollback
20 end

```

- (b) Update of StudentNo in CourseTaken
- (c) Delete from CourseTaken

```

1  -- For b) and c)
2  create trigger PhDStudMinOneCourse_PhDStud_Ins_Abort on CourseTaken after update,
   delete as
3  if exists (
4      select *
5      from Deleted D
6      where D.StudentNo not in (
7          select StudentNo
8          from CourseTaken ) )
9  begin
10     raiseerror ('Constraint Violation: A PhD student must take at least one course', 1,
11               1)
12     rollback
13 end

```

- (d) Delete from Course: Removing an entry from Course could indirectly affect the number of courses taken by one or several PhD students. This case, however, should be handled with the *on update cascade* option of the referential integrity constraint on the CourseNo field of CourseTaken.

3. A PhD student must take all courses taught by his/her supervisor.

This can be violated by:

- (a) Insert into PhDStudent
- (b) Update of Supervisor in PhDStudent

```

1  -- for a) and b)
2  -- Abort
3  create trigger StudAllCoursesOfSuperv_Stud_InsUpd_Abort on PhDStudent after insert,
   update as
4  if exists (
5      select *
6      from Inserted I
7      where exists (
8          select *
9          from Course C
10         where C.ProfNo = I.Supervisor
11             and C.CourseNo not in (
12                 select T.CourseNo
13                 from CourseTaken T
14                 where T.StudentNo = I.StudentNo ) ) )
15  begin
16     raiseerror ('Constraint Violation: A PhD student must take all the courses given by
17               his supervisor', 1, 1)
18     rollback
19 end
20 -- Repair
21 create trigger StudAllCoursesOfSuperv_Stud_InsUpd_Repair on PhDStudent after insert,
   update as
22 begin
23     insert into CourseTaken (StudentNo, CourseNo)
24         select I.StudentNo, C.CourseNo
25         from Inserted I, Professor P, Course C
26         where I.Supervisor = P.ProfNo
27             and C.ProfNo = P.ProfNo
28             and C.CourseNo not in (
29                 select T.CourseNo
30                 from CourseTaken T
31                 where T.StudentNo = I.StudentNo )
32 end

```

- (c) Insert into Course
- (d) Update of ProfNo in Course

```

1  -- Events c) and d)
2  -- Aborting the transaction, particularly in SQL Server (where triggers are executed
   immediately after the triggering instruction), would not work (well). The
   repairing rule being implicitly defined, or at least suggested, by the constraint
   (namely to automatically enrol the student in the added course), it will be the
   method of choice for this case.
3
4  -- Repair
5  create trigger StudAllCoursesOfSuperv_Course_InsUpd_Repair on Course after insert,
   update as
6  begin
7      insert into CourseTaken (StudentNo, CourseNo)
8      select S.StudentNo, I.CourseNo
9      from Inserted I, Professor P, PhDStudent S
10     where C.ProfNo = P.ProfNo
11           and S.Supervisor = P.ProfNo
12           and I.CourseNo not in (
13       select T.CourseNo
14       from CourseTaken T
15       where T.StudentNo = C.StudentNo )
16 end

```

(e) Update of StudentNo or CourseNo in CourseTaken

(f) Delete from CourseTaken

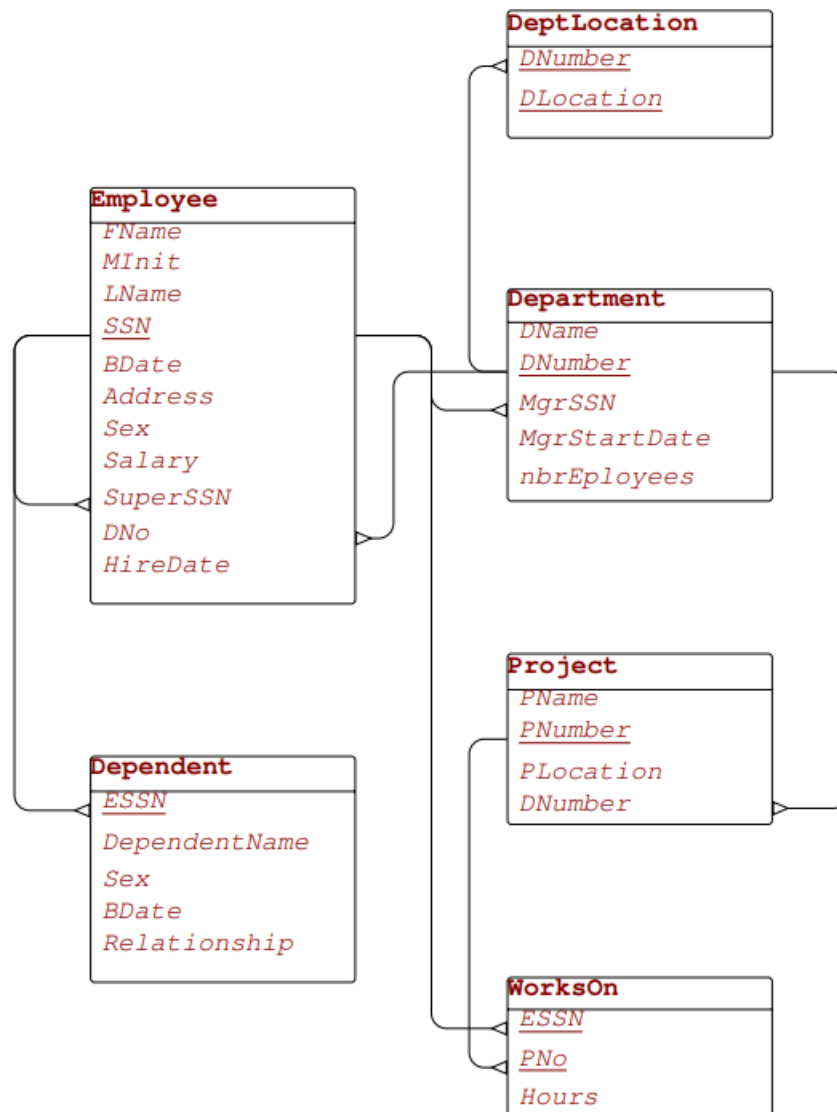
```

1  -- Events e) and f)
2  create trigger StudAllCoursesOfSuperv_CourseTaken_UpdDel_Abort on CourseTaken after
   update, delete as
3  if exists (
4      select *
5      from Deleted D, Course C, PhDStudent S
6      where D.CourseNo = C.CourseNo
7            and C.ProfNo = S.Supervisor
8            and D.StudentNo = S.StudentNo )
9  begin
10     raiserror ('Constraint Violation: A PhD student must take all the courses given by
   his supervisor', 1, 1)
11 end

```

2 Sessions 2 and 3: Exercises 'Employees department projects'

We have the schema:



In SQL Server, enforce the following constraints using a set of CHECK constraints, referential integrity constraints, or triggers.

Exercise 1. The age of employees must be greater than 18.

```

1 --Using a CHECK constraint
2 alter table Employee
3     add constraint employee_Age18
4     check ( dateadd(year,18,BDate) <= getdate() )
5
6 --Using a trigger
7 create trigger age18 on Employee after insert, update as
8 if exists (
9     select *
10    from Inserted
11   where dateadd(year,18,BDate) > getdate() )
12 begin
13     raiserror('Constraint Violation: The age of an employee must be greater than 18', 1, 1)
14     rollback
15 end
  
```

Exercise 2. The supervisor of an employee must be older than the employee.

```

1 --Using a trigger
  
```



```

2 create trigger supervisorAge on Employee after insert, update as
3 if exists (
4     select *
5     from Inserted I, Employee E
6     where ( I.SuperSSN = E.SSN and I.BDate < E.BDate )
7           or ( E.SuperSSN = I.SSN and E.BDate < I.BDate ) )
8 begin
9     raiserror( 'Constraint Violation: The age of an employee must be less than the age of his/
10             her supervisor', 1, 1)
11     rollback
12 end

```

Exercise 3. The salary of an employee cannot be greater than the salary of his/her supervisor.

```

1 -- Using a trigger
2 create trigger supervisorSalary on Employee after insert, update as
3 if exists (
4     select *
5     from Inserted I, Employee E
6     where ( I.SuperSSN = E.SSN and I.Salary > E.Salary )
7           or ( E.SuperSSN = I.SSN and E.Salary > I.Salary ) )
8 begin
9     raiserror('Constraint Violation: The salary of an employee cannot be greater than the salary
10             of his/her supervisor', 1, 1)
11     rollback
12 end

```

Exercise 4. The manager of a department must be an employee of that department.

```

1 -- Using UNIQUE and foreign key constraints
2 alter table Employee
3     add constraint UN_Employee_SSN_DNo unique( SSN, DNO )
4
5 alter table Department
6     add constraint FK_Employee_SSN_DNo foreign key( MgrSSN, DNumber ) references Employee( SSN,
7     DNo )

```

Exercise 5. The location of a project must be one of the locations of its department.

```

1 -- Using a foreign key constraint
2 alter table Project
3     add constraint FK_Project_DeptLocations foreign key( DNumber, PLocation ) references
4     DeptLocations( DNumber, DLocation )

```

Exercise 6. The hire date of employees must be greater than their birth date.

```

1 -- Using a CHECK key constraint
2 alter table Employee
3     add constraint HireDate_BDate check( HireDate > BDate )

```

Exercise 7. A supervisor must be hired at least 1 year before every employee s/he supervises.

```

1 -- Using a trigger
2 create trigger hireSuperv on Employee after insert, update as
3 if exists (
4     select *
5     from Inserted I, Employee E
6     where ( I.SuperSSN = E.SSN and datediff(year, E.HireDate, I.HireDate) < 1)
7           or ( E.SuperSSN = I.SSN and datediff(year, I.HireDate, E.HireDate) < 1))
8 begin
9     raiserror('Constraint Violation: A supervisor must be hired at least 1 year before every
10             employee s/he supervises', 1, 1)
11     rollback
12 end

```

Exercise 8. The attribute Department.NbrEmployees is a derived attribute from Employee.DNo.

```

1 -- Using value deriving triggers
2 create trigger DeptNbrEmp_Employee_InsUpdDel_Derive on Employee after insert, update, delete
3 as
4 begin

```

```

4  update Department D
5  set NbrEmployees = (
6      select count(*)
7      from Employee E
8      where E.DNo = D.DNumber )
9  where D.DNumber in (
10     select distinct I.DNo
11     from Inserted I )
12  or D.DNumber in (
13     select distinct D.DNo
14     from Deleted D )
15 end
16
17 -- Incremental version
18 create trigger derived_Department_NbrEmployees_Employee on Employee after insert, update,
19     delete as
20 begin
21     update Department
22     set NbrEmployees = NbrEmployees +
23         ( select count(*) from Inserted I where DNumber=I.DNo ) -
24         ( select count(*) from Deleted D where DNumber=D.DNo )
25     where DNumber in ( select DNo from Inserted )
26     or DNumber in ( select DNo from Deleted )
27 end
28 create trigger derived_Department_NbrEmployees_Department on Department after insert, update
29     as
30 if exists (
31     select *
32     from Inserted
33     where NbrEmployees <> (
34         select count(*)
35         from Employee E
36         where E.DNo = DNumber ) )
37 begin
38     raiserror('Constraint Violation: The attribute Department.NbrEmployees is a derived
39         attribute from Employee.DNo', 1, 1)
40     rollback
41 end

```

Exercise 9. An employee works at most in 4 projects.

```

1  -- Using a trigger
2  create trigger empNbrProj on WorksOn after insert, update as
3  if exists (
4      select *
5      from WorksOn W
6      group by W.ESSN
7      having count(*) > 4 )
8  begin
9      raiserror('Constraint Violation: An employee works at most in 4 projects', 1, 1)
10     rollback
11 end

```

Exercise 10. An employee works at least 30h/week and at most 50 h/week on all its projects.

```

1  -- Using a trigger
2  create trigger workson_30_50 on WorksOn after insert, update, delete as
3  if exists (
4      select *
5      from WorksOn
6      group by ESSN
7      having ( sum(Hours) < 30 )
8      or ( sum(Hours) > 50 ) )
9  begin
10     raiserror('Constraint Violation: An employee works at least 30 h/week and at most 50 h/week
11         on all its projects', 1, 1)
12     rollback
13 end

```

Exercise 11. Among all employees working on a project, at most 2 can work for less than 10 hours.

```

1  -- Using a trigger
2  create trigger worksonLess10h on WorksOn after insert, update as
3  if exists (
4      select *
5      from WorksOn
6      where Hours < 10
7      group by PNo
8      having count(*) > 2 )
9  begin
10     raiserror('Constraint Violation: A project can have at most 2 employees working on the
11         project less than 10 hours', 1, 1)
12     rollback
13 end

```

Exercise 12. Only department managers can work less than 5 hours on a project.

```

1  -- Using a set of triggers
2  create trigger worksonLess5h_WorksOn on WorksOn after insert, update as
3  if exists (
4      select *
5      from Inserted
6      where Hours < 5
7          and ESSN not in (
8              select MgrSSN
9              from Department
10             where MgrSSN is not null ) )
11  begin
12     raiserror('Constraint Violation: Only department managers can work less than 5 hours on a
13         project', 1, 1)
14     rollback
15 end
16
17 create trigger worksonLess5h_Department on Department after update, delete as
18 if exists (
19     select *
20     from Deleted
21     where MgrSSN not in (
22         select MgrSSN
23         from Department )
24     and MgrSSN in (
25         select ESSN
26         from WorksOn
27         where Hours < 5 ) )
28  begin
29     raiserror('Constraint Violation: Only department managers can work less than 5 hours on a
30         project', 1, 1)
31     rollback
32 end

```

Exercise 13. Employees that are not supervisors must work at least 10 hours on every project they work.

```

1  -- Using a set of triggers
2  create trigger workson10h_WorksOn on WorksOn after insert, update as
3  if exists (
4      select *
5      from Inserted
6      where Hours < 10
7          and ESSN not in (
8              select SuperSSN
9              from Employee
10             where SuperSSN is not null ) )
11  begin
12     raiserror('Constraint Violation: Employees that are not supervisors must work at least 10
13         hours on every project they work', 1, 1)
14     rollback
15 end
16
17 create trigger workson10h_Employee on Employee after update, delete as
18 if exists (
19     select *

```

```

19 from Deleted
20 where SuperSSN not in (
21     select SuperSSN
22     from Employee
23     where SuperSSN is not null )
24 and SuperSSN in (
25     select ESSN
26     from WorksOn
27     where Hours < 10 ) )
28 begin
29     raiseerror('Constraint Violation: Employees that are not supervisors must work at least 10
30         hours on every project they work', 1, 1)
31     rollback
32 end

```

Exercise 14. The manager of a department must work at least 5 hours on all projects controlled by the department.

```

1  -- Using a set of triggers
2  create trigger mgrProj_Department on Department after insert, update as
3  if exists (
4      select *
5      from ( Inserted I join Project P on I.DNumber = P.DNumber )
6          left outer join WorksOn on MgrSSN = ESSN and PNumber = PNo
7      where Hours is null
8          or Hours < 5 )
9  begin
10     raiseerror('Constraint Violation: A manager must work at least 5 hours on all projects
11         controlled by his/her department', 1, 1)
12     rollback
13 end
14
15 create trigger mgrProj_Project on Project after insert, update as
16 if exists (
17     select *
18     from ( Project P join Department D on D.DNumber = P.DNumber )
19         left outer join WorksOn on MgrSSN = ESSN and PNumber = PNo
20     where P.PNumber in (
21         select PNumber
22         from Inserted )
23         and ( Hours is null
24             or Hours < 5 ) )
25 begin
26     raiseerror('Constraint Violation: A manager must work at least 5 hours on all projects
27         controlled by his/her department', 1, 1)
28     rollback
29 end
30
31 create trigger mgrProj_WorksOn on WorksOn after update, delete as
32 if exists (
33     select *
34     from ( Department D join Project P on D.DNumber=P.DNumber)
35         left outer join WorksOn on MgrSSN = ESSN and PNumber = PNo
36     where D.MgrSSN in (
37         select ESSN
38         from Deleted )
39         and ( Hours is null
40             or Hours < 5 ) )
41 begin
42     raiseerror('Constraint Violation: A manager must work at least 5 hours on all projects
43         controlled by his/her department', 1, 1)
44     rollback
45 end

```

Exercise 15. The attribute Employee.SuperSSN is a derived attribute computed as follows. Department managers are supervised by the manager of Department 1 (Headquarters). Employees that are not managers are supervised by the manager of their department. Finally, the manager of Department 1 has a null value in attribute SuperSSN.

```

1  -- Using a set of triggers

```

```

2 create trigger derived_Employee_SuperSSN_Department on Department after insert, update as
3 if update(MgrSSN)
4 begin
5     update Employee
6     set SuperSSN = (
7         select case when SSN != D.MgrSSN
8             then D.MgrSSN
9             when SSN = D.MgrSSN and DNo != 1
10                then (
11                    select MgrSSN
12                    from Department
13                    where DNumber = 1 )
14                else null
15            end
16     from Department D
17     where DNo = D.DNumber )
18 -- if the department manager changes all employees of the department
19 -- must be updated
20     where ( DNo in (
21         select DNumber
22         from Inserted ) )
23 -- if the manager of department 1 changes, all department managers
24 -- must be updated
25     or ( 1 in (
26         select DNumber
27         from Inserted )
28     and SSN in (
29         select MgrSSN
30         from Department ) )
31 end
32
33 create trigger derived_Employee_SuperSSN_Employee on Employee after insert, update as
34 if update(DNo)
35 begin
36     update Employee
37     set SuperSSN = (
38         select case when SSN != MgrSSN
39             then D.MgrSSN
40             when SSN = MgrSSN and I.DNo != 1
41                then (
42                    select MgrSSN
43                    from Department
44                    where DNumber = 1 )
45                else null
46            end
47     from Inserted I, Department D
48     where SSN = I.SSN and I.DNo = D.DNumber )
49     where SSN in (
50         select SSN
51         from Inserted )
52 end

```

Exercise 16. The supervision relationship defined by Employee.SuperSSN must not be cyclic. (It is supposed that attribute Employee.SuperSSN is not derived as stated above.)

```

1 -- Using a trigger
2 create trigger noncyclic_subordinates on Employee after insert, update as
3 begin
4     create table #Supervision (
5         SSN char(9),
6         SuperSSN char(9)
7         primary key (SSN, SuperSSN) )
8
9     insert into #Supervision
10     select SSN, SuperSSN
11     from Employee
12     where SuperSSN is not null
13
14     while @@rowcount != 0 -- while previous operation affected some rows
15     begin
16         if exists (

```

```
17  select *
18  from #Supervision
19  where SSN = SuperSSN )
20  begin
21    raiserror('Constraint Violation: The supervision relationship is cyclic', 1, 1)
22    rollback
23  end
24
25  insert into #Supervision
26  select distinct S1.SSN, S2.SuperSSN
27  from #Supervision S1 join #Supervision S2 on S1.SuperSSN = S2.SSN
28  where not exists (
29    select *
30    from #Supervision S
31    where S.SSN = S1.SSN and S.SuperSSN = S2.SuperSSN )
32  end
33  end
```

Part II

Graph databases

3 Session 4: Relational graphs

Google released, in 2002, a subset of the structure of the WWW. In this dataset, web pages are represented by graph nodes such that when a web page A contains a hyperlink to web page B, a directed edge is created from node A to node B.

In this activity, we will focus on the performance of different queries. Therefore, we will use three PostgreSQL tables, which are subsets of different sizes of the web structure released by Google:

- webgraph1 table: 605 nodes (web pages) and 1521 edges (hyperlinks)
- webgraph2 table: 1622 nodes (web pages) and 6288 edges (hyperlinks)
- webgraph3 table: 4122 nodes (web pages) and 14356 edges (hyperlinks)

Below, we show the list of different uses cases we want to analyze:

Exercise 3.1. For each pair of connected nodes, find the 1-hop paths. Include four columns in the resultset: fromNode, toNode, length, path, which correspond to the source node, target node, length of the path, and visited nodes, respectively. Exclude repeated nodes in the path. That is, if there is an edges A -> A, do not consider that A, A, 1, A-A is a valid 1-path from A to A.

```
1 SELECT fromNode, toNode, 1 as length, fromNode || '-' || toNode as Path
2 FROM webgraph1
3 WHERE fromNode <> toNode;
```

Exercise 3.2. For each pair of connected nodes, find the 2-hop paths. Include four columns in the resultset: fromNode, toNode, length, path, which correspond to the source node, target node, length of the path and the visited nodes, respectively. Exclude repeated nodes in the path. That is, if A -> B and B->A are edges in the graph, do not consider that A, A, 2, A-B-A is a valid 2- path from A to A.

```
1 SELECT w1.fromNode, w2.toNode, 2 as length, w1.fromNode || '-' || w1.toNode || '-' || w2.
   toNode as Path
2 FROM webgraph1 w1
3 JOIN webgraph1 w2 ON w1.toNode = w2.fromNode
4 WHERE w1.fromNode <> w1.toNode AND w1.fromNode <> w2.toNode AND w2.fromNode <> w2.toNode;
```

Exercise 3.3. For each pair of connected nodes, find the 3-hop paths. Include four columns in the resultset: fromNode, toNode, length, path, which correspond to the source node, target node, length of the path and the visited nodes, respectively. Exclude repeated nodes in the path. That is, if A -> B, A->C and B->A, are edges in the graph, do not consider that A, C, 3, A-B-A-C is a valid 3-path from A to C.

```
1 SELECT w1.fromNode, w2.toNode, 2 as length, w1.fromNode || '-' || w1.toNode || '-' || w2.
   toNode as Path
2 FROM webgraph1 w1
3 JOIN webgraph1 w2 ON w1.toNode = w2.fromNode
4 JOIN webgraph1 w3 ON w2.toNode = w3.fromNode
5 WHERE w1.fromNode <> w3.toNode AND w1.fromNode <> w2.toNode AND w2.fromNode <> w3.toNode;
```

Exercise 3.4. for each pair of connected nodes, find the N-hop paths (the value of N is not known in advance). Include four columns in the result set: fromNode, toNode, length, path, which correspond to the source node, target node, length of the path and the visited nodes, respectively. Exclude repeated nodes in the path like in case "C".

```

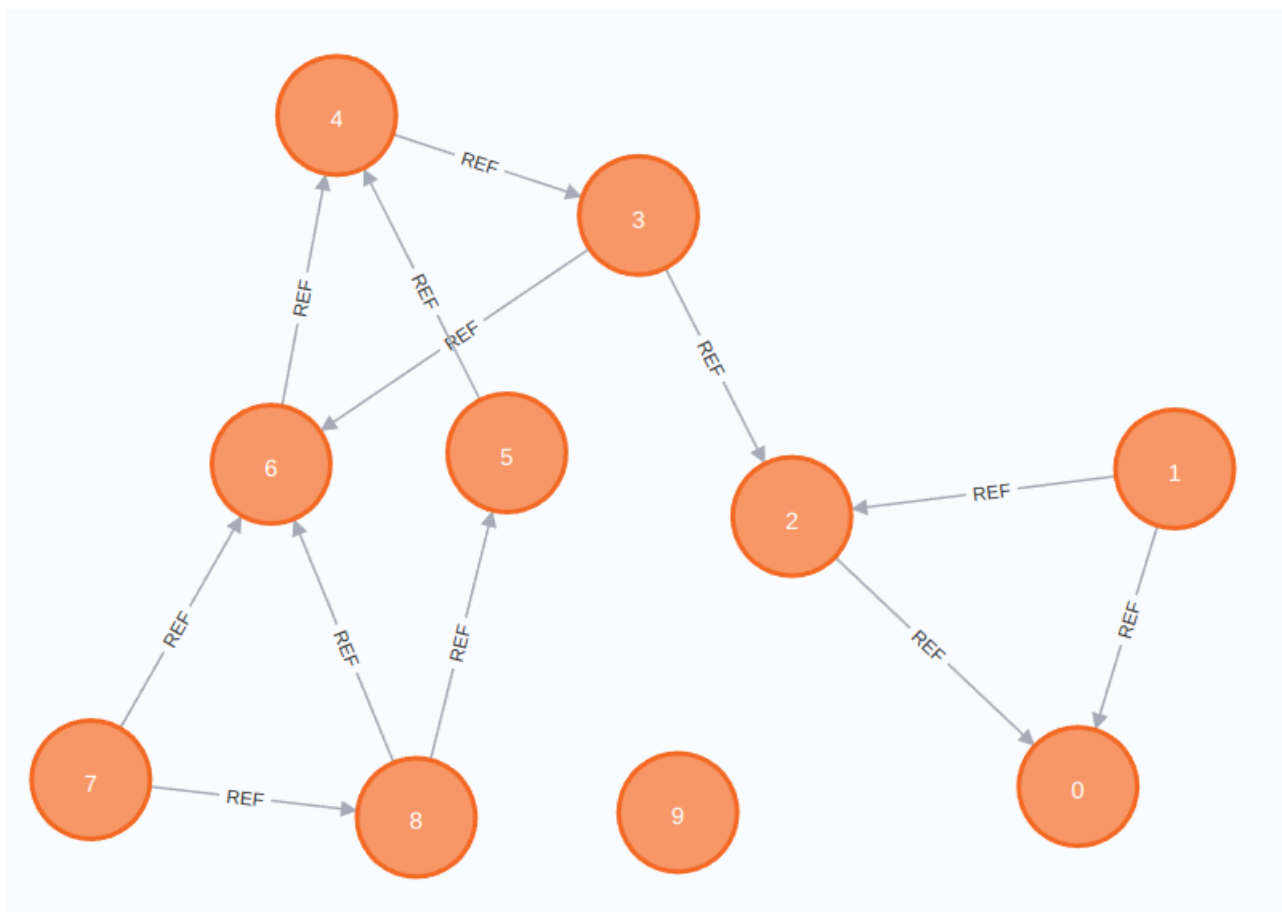
1 with recursive auxi(fromnode, tonode, long, path) as (
2   select wg.fromnode, wg.tonode, 1, '-' || wg.fromnode || '-' || wg.tonode
3   from webgraph1 wg
4   where wg.fromnode <> wg.tonode
5
6   union
7
8   select auxi.fromnode, wg.tonode, 1 + long, path || '-' || wg.tonode
9   from auxi, webgraph1 wg
10  where auxi.tonode = wg.fromnode
11        AND wg.fromnode <> wg.tonode
12        AND position('-', wg.tonode || '-' in auxi.path) = 0)
13
14 select * from auxi

```

4 Sessions 5 and 6: Neo4j and Cypher

4.1 Session 5: Cypher Part I

We will express queries over graph databases using Cypher, the high-level query language for Neo4J. To be able to easily check the correctness of the results, we will start with a small graph representing a subset of the web structure, as shown below:



Exercise 4.1. Vertex Degree

For each vertex compute its in-degree and out-degree. The result set is a list of vertices and both values. Those vertices that do not have outgoing edges and/or incoming edges, must not appear in the answer. (Note that the result set is not a graph).

I have solved this in two equivalent ways:


```

1 --1--
2 MATCH (p:URL)
3 WHERE size((:URL)-->(p))+size((p)-->(:URL))>0
4 RETURN p.name as url, size((:URL)-->(p)) as in_degree, size((p)-->(:URL)) as out_degree
5
6 --2--
7 MATCH (p:URL)
8 WITH p, size((:URL)-->(p)) AS in_degree, size((p)-->(:URL)) AS out_degree
9 WHERE in_degree+out_degree>0
10 RETURN p.name as url, in_degree, out_degree

```

Exercise 4.2. Vertex Degree variation

For each vertex calculate its in-degree and out-degree. The result set is a list of all vertices in the graph, together with the two values above, for each vertex. Those vertices that do not have outgoing edges and/or incoming edges, must appear in the answer with value “0”.

The way I solve the first exercise makes this question trivial:

```

1 MATCH (p:URL)
2 RETURN p.name as url, size((:URL)-->(p)) as in_degree, size((p)-->(:URL)) as out_degree
3
4 -- Another
5 MATCH (n)
6 OPTIONAL MATCH ()-[r1]->(n)
7 OPTIONAL MATCH (n)-[r2]->()
8 RETURN n.name, COUNT(distinct r1) as indegree, COUNT (distinct r2) as outdegree
9 ORDER BY n.name

```

Exercise 4.3. Calculating a maximum value

Find the maximum vertex in-degree. (the result set is not a graph).

I have solved this also in two different ways. I prefer the second solution, because it matches all nodes with maximum in-degree and is independent of the order done by the order by.

```

1 --1--
2 MATCH (p:URL)
3 RETURN p.name as url, size((:URL)-->(p)) as in_degree
4 ORDER BY in_degree DESC LIMIT 1
5
6 --2--
7 MATCH (p:URL)
8 WITH size((:URL)-->(p)) as in_degree
9 WITH max(in_degree) as mx
10 MATCH (q:URL)
11 WHERE size((:URL)-->(q)) = mx
12 RETURN q.name as url, mx
13
14 --3--
15 MATCH (n)
16 OPTIONAL MATCH (src)-->(n)
17 WITH n, COUNT(distinct src) as indegree
18 RETURN MAX(indegree)

```

Exercise 4.4. Find influential nodes

Find the subgraph which contains nodes whose in-degree is maximal in the graph (you should obtain only one node). Do the same for the out-degree (you should obtain four nodes).

For the in-degree:

```

1 MATCH (n:URL)
2 WITH size((:URL)-->(n)) as in_degree, n
3 WITH max(in_degree) as mx
4 MATCH (n2:URL)
5 WITH n2, size((:URL)-->(n2)) as in_degree
6 WHERE in_degree = mx
7 RETURN n2.name

```

And for the out-degree:

```

1 MATCH (n:URL)
2 WITH size((n)-->(:URL)) as out_degree, n
3 WITH max(out_degree) as mx
4 MATCH (n2:URL)
5 WITH n2, size((n2)-->(:URL)) as out_degree
6 WHERE out_degree = mx
7 RETURN n2.name

```

Exercise 4.5. Distance between nodes

For each pair of vertices, calculate the distance, i.e. the shortest simple path between them (without repeated edges in the path). Do not show the distance between two disconnected nodes (infinite distance). Exclude paths when source and target are the same node.

Although there is a specific function for this in Neo4j¹, we can manually do this:

```

1 MATCH path=(a:URL)-[*]->(b:URL)
2 WHERE a.name <> b.name
3 RETURN a.name, b.name, min(length(path)) as dist

```

Exercise 4.6. Distance between nodes using Cypher function

Solve the query in Use Case 5, but using the `shortestPath` built-in Cypher function. It has one parameter that represents a pattern path and returns the shortest path that matches this pattern. If there exists more than one shortest path, it returns any of them.

This one is an easier version of the previous exercise. Also, we see another way of computing the length of a path as `size(relationships(path))`.

```

1 MATCH (a:URL), (b:URL), sp=shortestpath((a)-[*]->(b))
2 WHERE a.name <> b.name
3 RETURN a.name, b.name, size(relationships(sp)) as dist
4
5 -- Another
6 MATCH path= shortestPath((n)-[*]->(p) )
7 WHERE n <> p
8 WITH nodes(path) as listanodos
9 RETURN [n in listanodos | n.name]

```

Exercise 4.7. Diameter

Compute the diameter of the graph, i.e. the longest distance between two nodes in the graph (excluding disconnected pairs of nodes).

This one's easy:

```

1 MATCH path=(:URL)-[*]->(:URL)
2 RETURN max(length(path)) as diameter
3
4 -- Another
5 MATCH p=(n1)-[*1..]->(n2)
6 WHERE n1 <> n2
7 WITH n1, n2, MIN(length(p)) as distance
8 RETURN MAX(distance)
9
10 -- Another
11 MATCH p=(n1)-[*1..]->(n2)
12 WHERE n1 <> n2
13 WITH n1.name, n2.name, MIN(length(p)) as distance
14 RETURN distance
15 ORDER BY distance DESC
16 LIMIT 1

```

¹See next exercise.

Exercise 4.8. webgraph3

Repeat use cases 1 to 7 using the webgraph3 database, which represents the same information as the corresponding relational database in Activity 1 (Session 4).

We just have to execute the same queries in the webgraph3 database.

Exercise 4.9. Paths

Compute all the 1, 2, 3, and n-hops in the graph, and compare against the results obtained using PostgreSQL in Activity 1. Note: start with “limit X”, increasing “X” to prevent that the algorithm runs indefinitely.

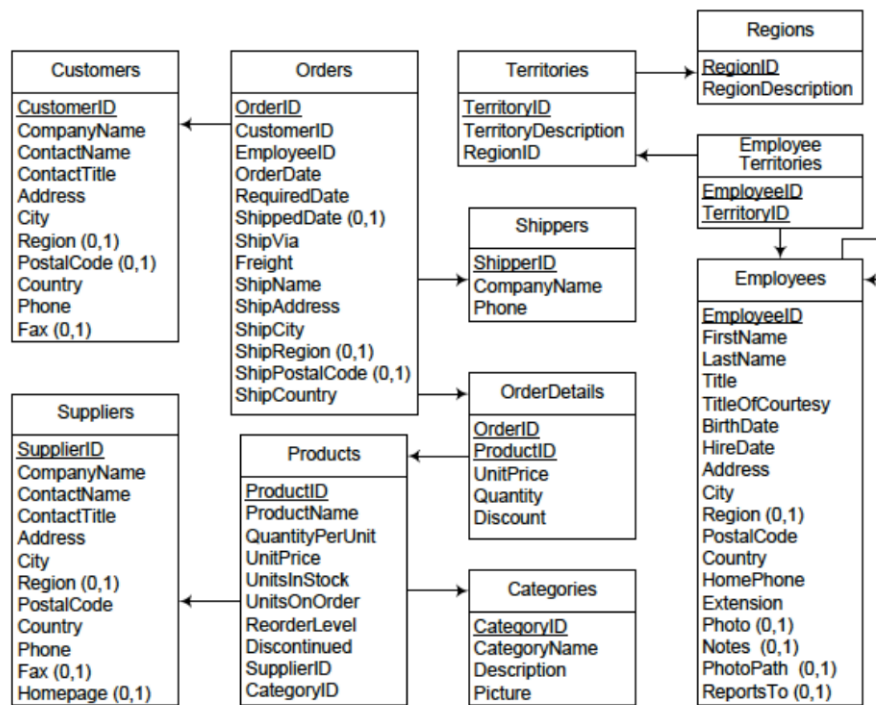
```

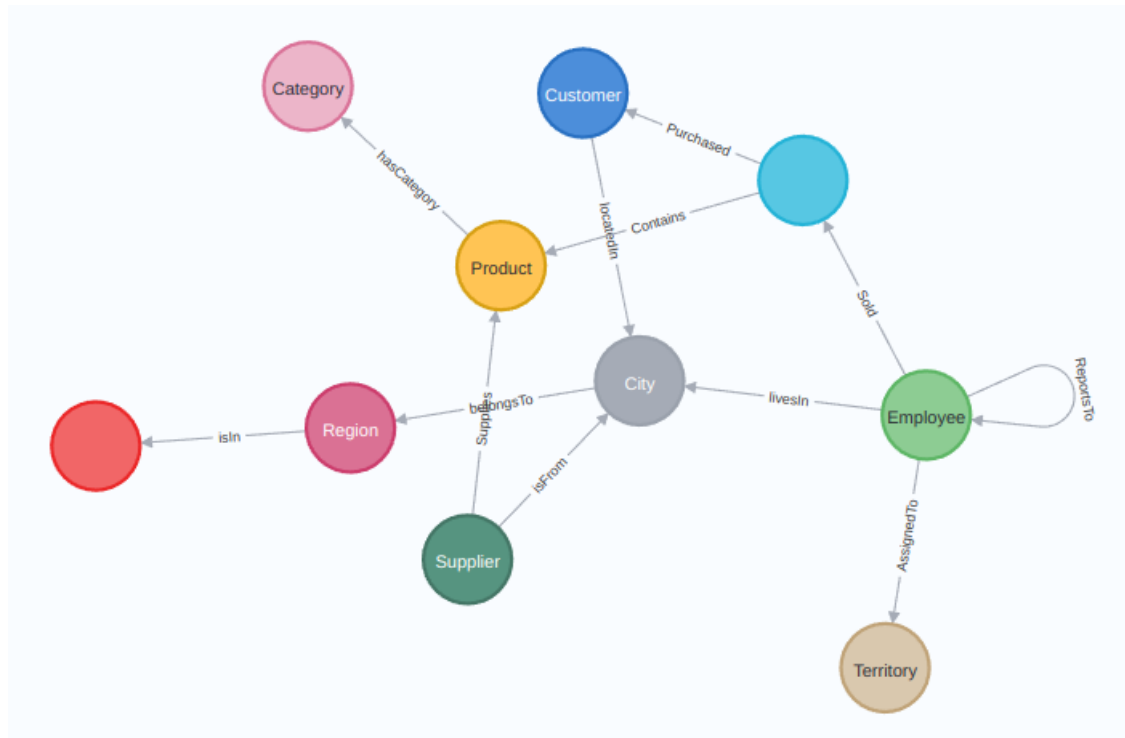
1 MATCH path = (end)<-[*1..]-(start)
2 WHERE ALL (n in nodes(path) where 1 = size([m in nodes(path) where m=n]))
3 RETURN start.name, LENGTH(path) AS length, [p in NODES(path) | p.name], end.name
4 ORDER BY length
5 MATCH path = (end)<-[*1..]-(start:URL{name:745315})
6 WHERE ALL (n in nodes(path) where
7 1 = size([m in nodes(path) where m=n]))
8 RETURN LENGTH(path) AS length, [p in NODES(path) | p.name]
9 ORDER BY length desc

```

4.2 Session 6: Cypher Part II

Exercise 4.10. Consider the Northwind database, whose schema is:





Write in Cypher the following queries over the northwindhg.db database:

1. List products and their unit price.

```
1 MATCH (v:Product)
2 RETURN v.productID, v.unitPrice
```

2. List information about products 'Chocolade' & 'Pavlova'.

```
1 MATCH (p:Product)
2 WHERE p.productName IN ['Chocolade', 'Pavlova']
3 RETURN p
```

3. List information about products with names starting with a "C", whose unit price is greater than 50.

```
1 MATCH (p:Product)
2 WHERE p.productName STARTS WITH "C" AND tofloat(p.unitPrice) > 50
3 RETURN p.productName, p.unitPrice
```

4. Same as 3, but considering the sales price, not the product's price.

```
1 MATCH (p:Product) <- [c:Contains] - (o:Order)
2 WHERE p.productName STARTS WITH "C" AND tofloat(c.unitPrice) > 50
3 RETURN distinct p.productName, p.unitPrice, c.unitPrice
```

5. Total amount purchased by customer and product.

```
1 MATCH (c:Customer)
2 OPTIONAL MATCH (p:Product) <- [pu:Contains] - (:Order) - [:Purchased] -> (c)
3 RETURN c.customerName, p.productName, tofloat(sum(pu.unitPrice) * pu.quantity) as volume
4 ORDER BY volume desc
```

6. Top ten employees, considering the number of orders sold.

```
1 MATCH (:Order) <- [:Sold] - (e:Employee)
2 RETURN e.firstName, e.lastName, count(*) AS Ordenes
3 ORDER BY Ordenes DESC LIMIT 10
```

7. For each employee, list the assigned territories.

```
1 MATCH (t:Territory)<-[:AssignedTo]-(e:Employee)
2 RETURN e.lastName, COLLECT(t.name)
```

8. For each city, list the companies settled in that city.

```
1 MATCH (c:City)<-[:locatedIn]-(c1:Customer)
2 RETURN c.cityname, COLLECT(c1.customerName)
```

9. How many persons an employee reports to, either directly or transitively?

```
1 MATCH (report:Employee)
2 OPTIONAL MATCH (e)<-[:ReportsTo*]-(report)
3 RETURN report.lastName AS e1, COUNT(rel) AS reports
```

10. To whom do persons called “Robert” report to?

```
1 MATCH (e:Employee)<-[:ReportsTo*]-(sub:Employee)
2 WHERE sub.firstName = 'Robert'
3 RETURN e.firstName, e.lastName, sub.lastName
```

11. Who does not report to anybody?

```
1 MATCH (e:Employee)
2 WHERE NOT (e)-[:ReportsTo]->()
3 RETURN e.firstName AS TopBossFirst, e.lastName AS TopBossLast
```

12. Suppliers, number of categories they supply, and a list of such categories

```
1 MATCH (s:Supplier)-->(:Product)-->(c:Category)
2 WITH s.supplierName AS Supplier, collect(distinct c.categoryName) AS Categories
3 WITH Supplier, Categories, size(Categories) AS Cantidad ORDER BY Cantidad DESC
4 RETURN Supplier, Cantidad, Categories
```

13. Suppliers who supply beverages

```
1 MATCH (c:Category {categoryName:"Beverages"})<--(:Product)<--(s:Supplier)
2 RETURN DISTINCT s.supplierName AS ProduceSuppliers
```

14. Customer who purchases the largest amount of beverages

```
1 MATCH (cust:Customer)<-[:Purchased]-(o:Order)-[o1:Contains]->(p:Product), (p)-[:hasCategory]
->
2 (c:Category{categoryName:"Beverages"})
3 RETURN cust.customerName AS CustomerName, SUM(o.quantity)
4 LIMIT 1
```

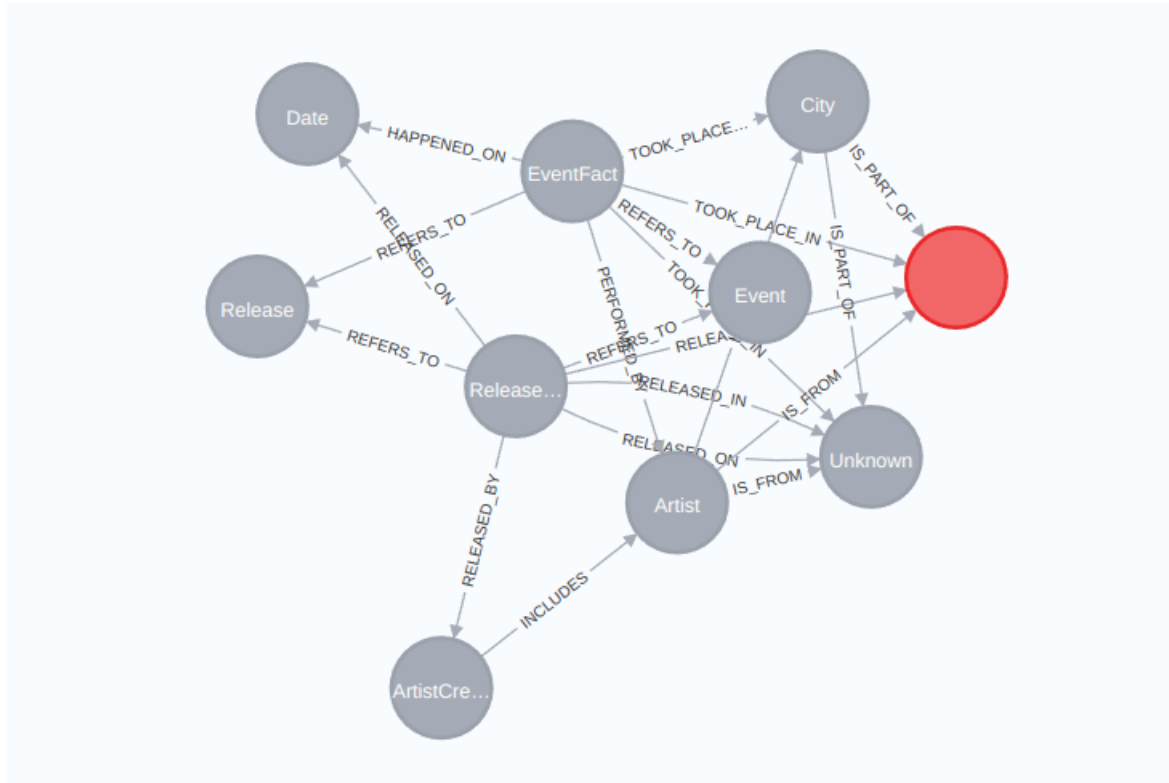
15. List the five most popular products (considering number of orders)

```
1 MATCH (c:Customer)<-[:Purchased]-(o:Order)-[o1:Contains]->(p:Product)
2 RETURN c.customerName, p.productName, count(o1) AS orders
3 ORDER BY orders desc LIMIT 5
```

16. Products ordered by customers from the same country than their suppliers

```
1 MATCH (c:Customer) -[r:locatedIn]->(cy:City) -[:belongsTo]->(:Region) -[:isIn]->(co:Country)
2 WITH co, c MATCH (s:Supplier) WHERE co.countryname = s.country
3 WITH s, co, c MATCH(s)-[su:Supplies]-(p:Product)<-[:Contains]-(o:Order)-[:Purchased]->(c)
4 RETURN c.customerName, s.supplierName, co.countryname, p.productName
```

Exercise 4.11. Switch to the MusicBrainz database, doing the same steps as in Assignment 2. Now, the database is musicbrainz. The schema is:



1. Compute the total number of releases per artist.

```
1 MATCH (r:ReleaseFact)-[rb:RELEASED_BY]->(ac:ArtistCredit)-[inc:INCLUDES]->(a:Artist)
2 RETURN a.name, COUNT(r)
```

2. Compute the total number of releases per artist and per year.

```
1 MATCH (d:Date)<-[ro:RELEASED_ON]-(r:ReleaseFact)-[rb:RELEASED_BY]->(ac:ArtistCredit)-[inc:INCLUDES]->(a:Artist)
2 RETURN a.name, d.year, COUNT(r)
3 ORDER BY a.name, d.year
```

3. Compute the total number of events per artist.

```
1 MATCH (e:Event)<-[rt:REFERS_TO]-(ef:EventFact)-[pb:PERFORMED_BY]->(a:Artist)
2 RETURN a.name, COUNT(e)
```

4. Compute the number of times the artist performed in each event.

```
1 MATCH (e:Event)<-[rt:REFERS_TO]-(ef:EventFact)-[pb:PERFORMED_BY]->(a:Artist)
2 RETURN a.name, e.name, COUNT(ef)
```

5. For each (event, artist, year) triple, compute the number of times the artist performed in an event on an year.

```
1 MATCH (e:Event)<-[rt:REFERS_TO]-(ef:EventFact)-[pb:PERFORMED_BY]->(a:Artist)
2 MATCH (d:Date)<-[ho:HAPPENED_ON]-(ef)
3 RETURN a.name, e.name, d.year, COUNT(ef)
```

6. Same as Query 5, for artists in the United Kingdom and events happened after year 2006.

```

1 MATCH (d:Date)<-[ho:HAPPENED_ON]-(ef)
2 WHERE d.year > 2006
3 MATCH (e:Event)<-[rt:REFERS_TO]-(ef:EventFact)-[pb:PERFORMED_BY]->(a:Artist)-[isf:IS_FROM]
  ]->(c:Country)
4 WHERE c.name = 'United Kingdom'
5 RETURN a.name, e.name, d.year, COUNT(ef)

```

7. Compute the number of releases, per language, in the UK.

```

1 MATCH (r:Release)<-[rt:REFERS_TO]-(rf:ReleaseFact)-[ri:RELEASED_IN]->(c:Country)
2 WHERE c.name='United Kingdom'
3 RETURN r.language, count(r)

```

8. Compute, for each pair of artists, the number of times they have performed together at least twice in an event.

```

1 MATCH (a1:Artist)<-[pb1:PERFORMED_BY]-(ef:EventFact)-[pb2:PERFORMED_BY]->(a2:Artist)
2 WHERE a1.name < a2.name
3 MATCH (e:Event)<-[rt:REFERS_TO]-(ef)
4 WITH a1.name as artist1, a2.name as artist2, COUNT(e) as times
5 WHERE times>=2
6 RETURN artist1, artist2, times

```

9. Compute the triples of artists, and the number of times they have performed together in an event, if this number is at least 3.

```

1 MATCH (a1:Artist)<-[pb1:PERFORMED_BY]-(ef:EventFact)-[pb2:PERFORMED_BY]->(a2:Artist)
2 MATCH (ef)-[pb3:PERFORMED_BY]->(a3:Artist)
3 WHERE a1.name < a2.name < a3.name
4 MATCH (e:Event)<-[rt:REFERS_TO]-(ef)
5 WITH a1.name as artist1, a2.name as artist2, a3.name as artist3, COUNT(e) as times
6 WHERE times>=3
7 RETURN artist1, artist2, artist3, times

```

10. Compute the quadruples of artists, and the number of times they have performed together in an event, if this number is at least 3.

```

1 MATCH (a1:Artist)<-[pb1:PERFORMED_BY]-(ef:EventFact)-[pb2:PERFORMED_BY]->(a2:Artist)
2 MATCH (a4:Artist)<-[pb4:PERFORMED_BY]-(ef)-[pb3:PERFORMED_BY]->(a3:Artist)
3 WHERE a1.name < a2.name < a3.name < a4.name
4 MATCH (e:Event)<-[rt:REFERS_TO]-(ef)
5 WITH a1.name as artist1, a2.name as artist2, a3.name as artist3, a4.name as artist4,
  COUNT(e) as times
6 WHERE times>=3
7 RETURN artist1, artist2, artist3, artist4, times

```

11. Compute the pairs of artists that have performed together in at least two events and that have worked together in at least one release, returning the number of events and releases together.

```

1 MATCH (a1:Artist)<-[pb1:PERFORMED_BY]-(ef:EventFact)-[pb2:PERFORMED_BY]->(a2:Artist)
2 WHERE a1.name < a2.name
3 MATCH (e:Event)<-[rte:REFERS_TO]-(ef)
4 WITH a1, a2, COUNT(e) as events
5 MATCH (a1)<-[in1:INCLUDES]-(ac:ArtistCredit)-[in2:INCLUDES]->(a2)
6 MATCH (ac)<-[rb:RELEASED_BY]-(rf:ReleaseFact)-[rtr:REFERS_TO]->(r:Release)
7 WITH a1.name as artist1, a2.name as artist2, events, COUNT(r) as releases
8 WHERE events>=2 AND releases>=1
9 RETURN artist1, artist2, events, releases

```

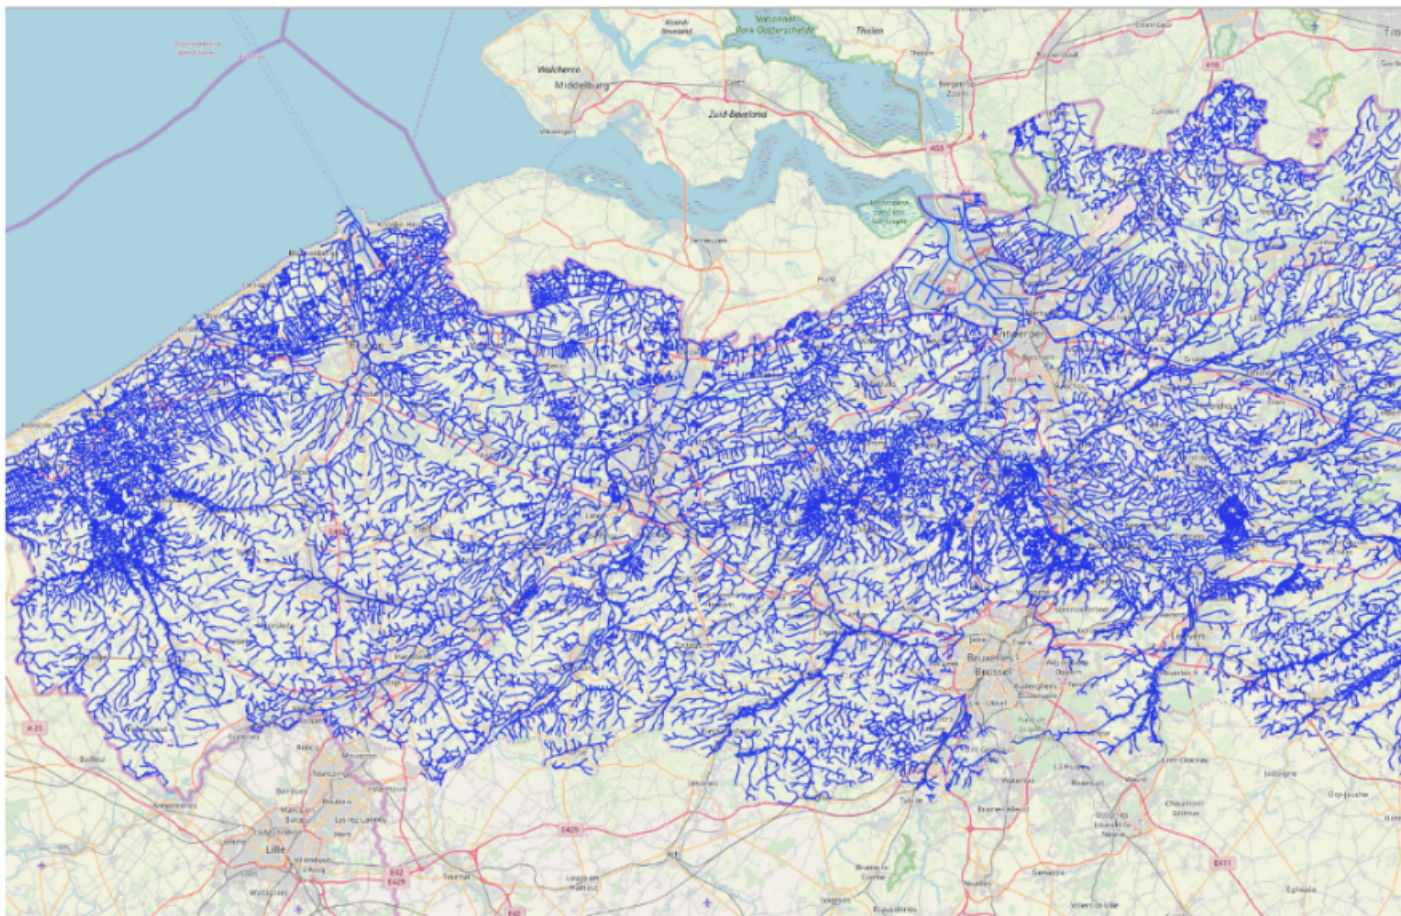
12. Compute the number of artists who released a record and performed in at least an event, and the year(s) this happened.


```

1 MATCH (a1:Artist)<-[pb1:PERFORMED_BY]-(ef:EventFact)-[ho:HAPPENED_ON]->(de:Date)
2 MATCH (e:Event)<-[rte:REFERS_TO]-(ef)
3 WITH a1, COUNT(e) as events, COLLECT(DISTINCT(de.year)) as eventyears
4 MATCH (a1)<-[in1:INCLUDES]-(ac:ArtistCredit)<-[rb:RELEASED_BY]-(rf:ReleaseFact)
5 MATCH (dr:Date)<-[ro:RELEASED_ON]-(rf)-[rtr:REFERS_TO]->(r:Release)
6 WITH a1.name as artist1, events, eventyears, COUNT(r) as releases, COLLECT(DISTINCT(dr.
   year)) as releaseyears
7 WHERE events>=1 AND releases>=1
8 RETURN artist1, events, releases, eventyears, releaseyears

```

Exercise 4.12. We will query the Flanders river system depicted in Figure 1. The schema and properties are shown in the next figures. Segments are represented as nodes, with label :Segment (and their corresponding properties), and the relation between the nodes is called :flowsTo, defined as follows: there is a relation :flowsTo from node A to node B if the water flows to segment B from segment A. This is stored in the rivers database.



1. Query 1. Compute the average segment length.

```

1 MATCH (n:Segment)
2 RETURN avg(n.length) AS avglength

```

2. Query 2. Compute the average segment length by segment category

```

1 MATCH (n:Segment)
2 RETURN n.catc as category, avg(n.length)
3 AS avglength order by category asc

```


3. Query 3. Find all segments that have a length within a 10% margin of the length of segment with ID 6020612.

```
1 MATCH (n:Segment {vhas:6020612})
2 WITH n.length AS length
3 MATCH (m:Segment)
4 WHERE m.length < length*1.1 AND m.length > length*0.9
5 RETURN m.vhas, m.length;
```

4. Query 4. For each segment find the number of incoming and outgoing segments.

```
1 MATCH (src:Segment)-[:flowsTo]->(n:Segment)-[:flowsTo]->(target:Segment)
2 RETURN n.vhas AS nodenbr, COUNT(DISTINCT src) AS segIn,
3 COUNT(DISTINCT target) AS segOut
```

5. Query 5. Find the segments with the maximum number of incoming segments.

```
1 MATCH (n:Segment)
2 OPTIONAL MATCH (src:Segment)-[:flowsTo]->(n)
3 WITH n, COUNT(DISTINCT src) AS indegree
4 WITH COLLECT([n, indegree]) AS tuples, MAX(indegree) AS max
5 RETURN [t IN tuples WHERE t[1] = max | t]
```

6. Query 6. Find the number of splits in the downstream path of segment 6020612.

```
1 MATCH (n:Segment {vhas:6020612})
2 CALL apoc.path.spanningTree(n,{relationshipFilter:"flowsTo>", minLevel: 1}) YIELD
3 path AS pp
4 UNWIND NODES(pp) AS p
5 MATCH (p)-[:flowsTo]->(r:Segment) WITH p, count(DISTINCT r) AS co WHERE co > 1
6 RETURN count(p)
```

7. Query 7: Find the number of in-flowing segments in the downstream path of segment 6020612.

```
1 MATCH (n:Segment {vhas:6020612})
2 CALL apoc.path.spanningTree(n,{relationshipFilter:"flowsTo>", minLevel: 1}) YIELD path AS
3 pp
4 WITH [p IN NODES(pp) | p.vhas] AS ids
5 UNWIND ids AS id
6 WITH collect(DISTINCT id) AS ids
7 MATCH (s:Segment)-[:flowsTo]->(p)
8 WHERE NOT s.vhas IN ids AND p.vhas <> 6020612 AND p.vhas IN ids
9 RETURN count(DISTINCT s) AS inflows
```

8. Query 8. Determine if there is a loop in the downstream path of segment 6031518.

```
1 MATCH (n:Segment {vhas:6031518})
2 CALL apoc.path.spanningTree(n, {relationshipFilter:"flowsTo>", minLevel: 1}) YIELD path
3 AS pp
4 WITH [p IN NODES(pp) | p] AS nodelist
5 UNWIND nodelist AS p
6 CALL apoc.path.expandConfig(p, {relationshipFilter:"flowsTo>", minLevel: 1,
7 terminatorNodes:[p], whitelistNodes:nodelist}) yield path AS loop
8 RETURN count(loop) > 0 AS loops
```

9. Query 9. Find the length, the # of segments, and the IDs of the segments, of the longest branch of upstream flow starting from a given segment.

```
1 MATCH (n:Segment {vhas:6020612})
2 CALL apoc.path.expandConfig(n,{relationshipFilter:"<flowsTo", minLevel: 1}) YIELD path AS
3 pp
4 WITH reduce(longi= tofloat(0), n IN nodes(pp)| longi+ tofloat(n.length)) AS blength,
5 Length(pp) AS alength, [p IN NODES(pp) | p.vhas] AS nodelist
6 WITH blength, alength, nodelist[size(nodelist)-1] AS id
7 WITH id, max(blength) AS ml, collect([id,blength,alength]) AS coll
8 UNWIND lhops AS hops
9 RETURN id,ml,hops ORDER BY id DESC;
```

10. Query 10. How many paths exist between two given segments X and Y?

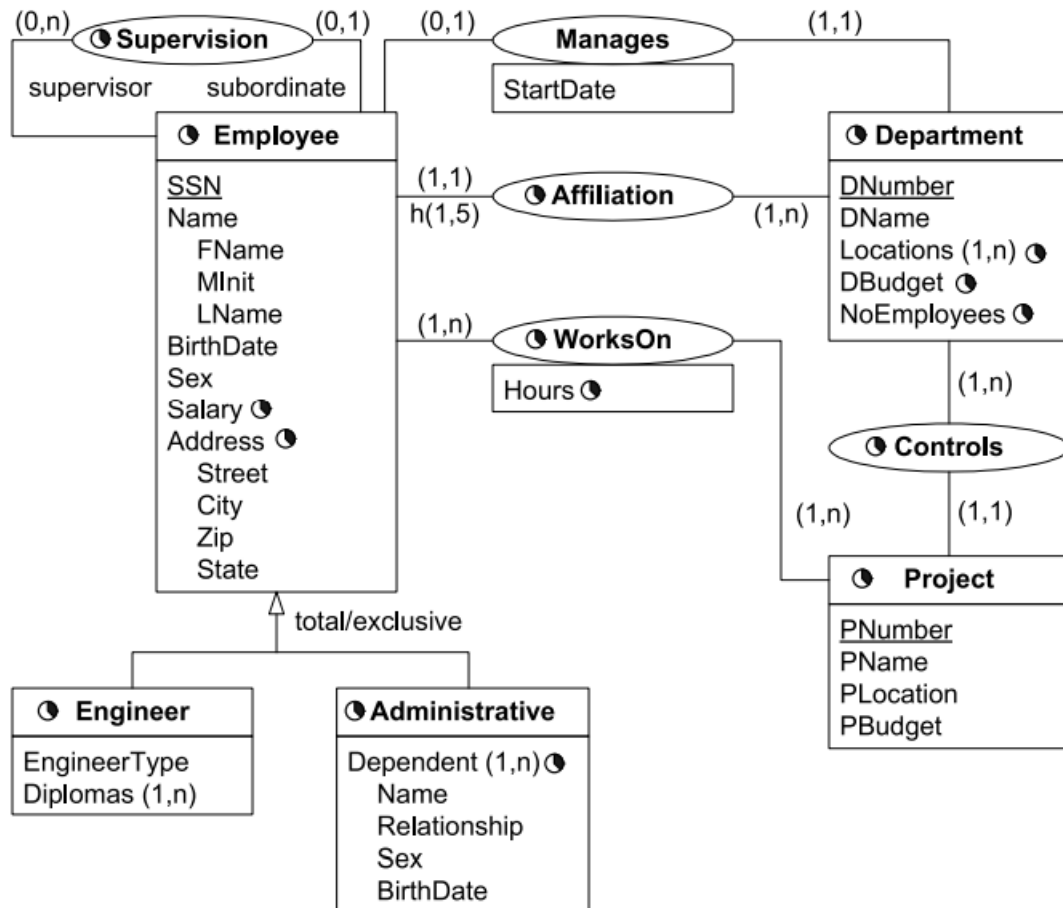
```
1 MATCH (n:Segment {vhas:6020612}),(m:Segment {vhas: 7036554})
2 CALL apoc.path.expandConfig(n,{relationshipFilter:"<flowsTo", minLevel: 1,
   terminatorNodes:[m]}) yield path as pp
3 RETURN count(pp) as paths
```

11. Query 11. Find all segments reachable from the segment closest to Antwerpen's Groenplaats

```
1 CALL apoc.spatial.geocodeOnce('Groenplaats Antwerpen Flanders Belgium') YIELD location as
   ini
2 MATCH (n:Segment)
3 WITH n, ini,distance(point({longitude:n.source_long, latitude:n.source_lat}), point({
   longitude:ini.longitude, latitude:ini.latitude})) as d
4 WITH n, d order by d asc limit 1
5 CALL apoc.path.spanningTree(n,{relationshipFilter:"flowsTo>", minLevel: 1}) YIELD path as
   pp
6 UNWIND NODES(pp) as p
7 RETURN p.vhas;
```

Part III

Temporal databases



Part IV

Spatial databases

5 Lab 10

In the directory **tp10**, perform:

```

1 // Initialization
2 $ createdb infoh415
3 $ psql infoh415 -c "CREATE EXTENSION postgis;"
4
5 // Tables creation
6 $ psql infoh415
7 infoh415$ \i create_tables.sql
8
9 // Tables population
10 infoh415$ \i insertion.sql
11 infoh415$ \i create_index.sql
12
13 // Importing and exporting
14 infoh415$ quit
15 $ cd shapefiles/
16 $ pgsql2shp infoh415 regions
17 $ pgsql2shp infoh415 cities
18 $ shp2pgsql -W "latin1" bel_city.shp > ../shp_insert.sql
19 $ shp2pgsql -W "latin1" bel_dist.shp >> ../shp_insert.sql
20 $ shp2pgsql -W "latin1" bel_prov.shp >> ../shp_insert.sql
21 $ shp2pgsql -W "latin1" bel_regn.shp >> ../shp_insert.sql
22 $ shp2pgsql -W "latin1" belriver.shp >> ../shp_insert.sql
23 $ cd ..
24 $ psql infoh415
25 infoh415$ \i shp_insert.sql
26
27 // Specify the SRID to use WGS84
28 infoh415$ \i update_srid.sql

```

Write down and execute the following queries:

1. Get the SRID from table cities.

```

1 select st_srid(geom) as SRID
2 from bel_city
3 limit 5;

```

2. Get a textual description for this SRID.

```

1 select srtext
2 from spatial_ref_sys
3 where srid=4326;

```

3. Get the dimension of geographical objects in that table.

```

1 select distinct st_dimension(geom) as dim, st_coorddim(geom) as coorddim
2 from bel_city;

```

4. Get the geometry type of these objects.

```

1 select distinct st_geometrytype(geom) as type
2 from bel_city;

```

5. Compute the distance between the cities of IXELLES and BRUGGE.

```

1 select st_distance(
2   st_transform((select geom from bel_city where name = 'Ixelles'),3812),
3   st_transform((select geom from bel_city where name = 'Brugge'),3812)) as dist,
4 st_distancesphere(
5   (select geom from bel_city where name = 'Ixelles'),
6   (select geom from bel_city where name = 'Brugge')) as dist_sphere

```

The number 3812 in the transform function refers to Belgian Lambert 2008, which is a conic representation adapted for Belgium. If we don't apply this transformation, we get a distance measured in degrees, which has little meaning for us.


6. Compute the bounding rectangle for the BRABANT province.

```

1 create table bbox as
2   select st_envelope(geom)
3   from bel_prov
4   where name = 'Brabant';

```

We create the table to be able to visualize the result in QGIS:



7. Compute the geographical union of the bel_reg and bel_prov tables.

```

1 create table geom_union as
2   select st_union(geom)
3   from (select geom from bel_regn
4         union
5         select geom from bel_prov)
6   as tmp;

```



8. Compute the length of each river.

```
1 select st_length2dspheroid(geom, 'SPHEROID["GRS_1980",6378137,298.257222101]')
2 from belriver;
```

The second argument is the spheroid that the function should use to compute the lengths, as its definition is *ST_Length2DSpheroid(geometry geom, spheroid sp)*.

9. Create a table containing all cities that stand less than 1000m from a river.

```
1 select name, dist
2 from
3   (select c.name as name, min(st_distancesphere(c.geom, r.geom)) as dist
4    from bel_city c, belriver r
5    group by c.name) as tmp
6 where dist < 1000;
7
8 -- or
9
10 select c.name as name, min(st_distancesphere(c.geom, r.geom)) as dist
11 from bel_city c
12 join belriver r on st_distancesphere(c.geom, r.geom) < 1000
13 group by c.name;
```

10. For each river, compute the length of its path inside each province it traverses.

```
1 select r.name as river, p.name as province,
2        st_length2dspheroid(st_intersection(r.geom, p.geom),
3                             'SPHEROID["GRS_1980",6378137,298.257222101]') as inside_len
4 from belriver r
5 join bel_prov p on st_intersects(r.geom, p.geom);
```

6 Lab 11

Download http://bioge.ucdavis.edu/data/climate/worldclim/1_4/grid/cur/alt_10m_bil.zip and extract it to **tp11**.

In the directory **tp11** perform:

```

1 createdb tp11
2 psql tp11 -c "CREATE EXTENSION postgis;"
3 psql tp11 -c "CREATE EXTENSION postgis_raster;"
4 psql tp11 -f generate.sql
5
6 cd bel_alt/
7 raster2pgsql BEL_alt.vrt > insert_bel_alt.sql
8 psql tp11 -f insert_bel_alt.sql
9 cd ../alt_10m_bil/
10 raster2pgsql alt.bil > insert_alt.sql
11 psql tp11 -f insert_alt.sql
12 cd ..
13
14
15 # 3.2 Update rasters
16 psql tp11
17 SELECT UpdateRasterSRID('alt', 'rast', 4326);
18 SELECT UpdateRasterSRID('bel_alt', 'rast', 4326);

```

Write down and execute the following queries:

1. Compute the difference between the two altitude datasets. Try to perform this exercise first without using ST_Resample. Export the result and visualize it in QGIS:

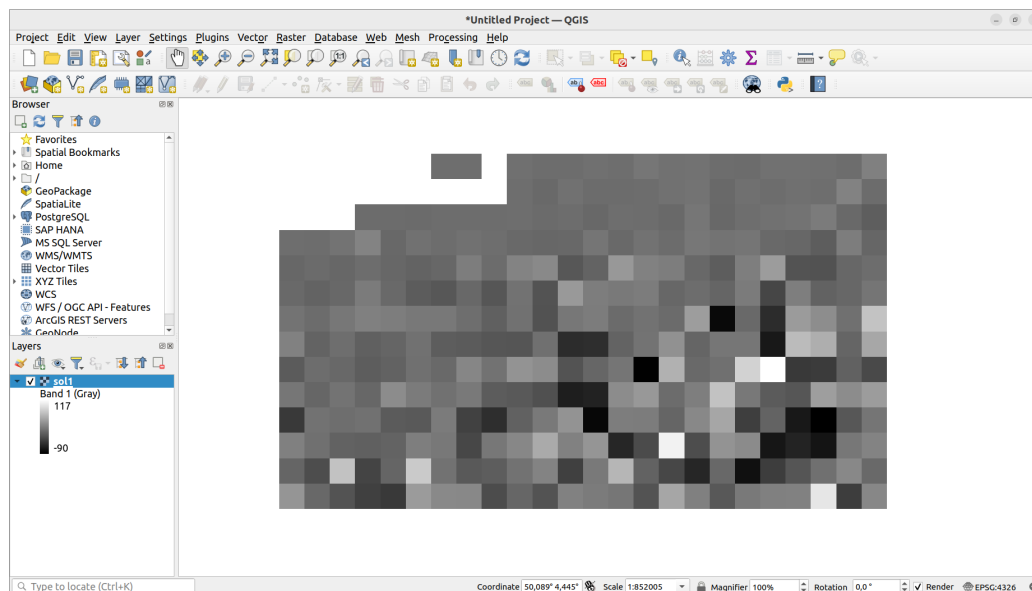
```
1 gdal_translate -of GTiff PG:"dbname=tp11 schema=public table=sol1" sol1.tiff
```

```

1 create table sol1 as
2 SELECT ST_MapAlgebra( a.rast, 1, b.rast, 1, '[rast2] - [rast1]') AS rast
3 FROM
4     (select b.RID, ST_Resample(b.rast, a.rast) as rast
5      from bel_alt b, alt a)
6 as b, alt a;

```

Once we have the table, we execute the command to export the raster and then we import it to QGIS. The result is (not very awesome):



2. Compute the maximum altitude in Belgium.

```

1 SELECT (ST_SummaryStats(rast)).max as max
2 FROM bel_alt;

```

3. Get the altitudes of all cities in Belgium.

```

1 SELECT name, ST_Value(rast, geom)
2 FROM bel_alt
3 JOIN bel_city ON ST_Intersects(geom, rast);

```

4. Compute the maximum and minimum altitudes for each province.

```

1 SELECT name, (stats).max, (stats).min
2 FROM (
3   SELECT name, ST_SummaryStats(ST_Clip(rast, 1, geom, TRUE)) AS stats
4   FROM alt
5   JOIN bel_prov ON ST_Intersects(geom, rast))
6 AS foo;

```

5. Create a new raster table restraining the alt_16 raster to Belgium. (Hint: use ST_Intersection.)

```

1 CREATE TABLE alt_16_belgium AS
2 SELECT rid, ST_Clip(rast, 1, geom, TRUE) AS rast
3 FROM (SELECT ST_UNION(ARRAY(SELECT geom FROM bel_regn)) AS geom) AS belgium
4 JOIN alt ON ST_Intersects(geom, rast);

```

6. Compute the altitude along each river.

```

1 SELECT name, ST_AsText((position).geom), ST_Value(rast, (position).geom)
2 FROM (SELECT name, rast, ST_DumpPoints(ST_Segmentize(geom, 0.1)) AS position
3      FROM belriver
4      JOIN bel_alt on ST_Intersects(geom, rast))
5 AS dp;

```

Spatial functions reference

- **ST_CLIP**: Returns a raster that is clipped by the input geometry geom. If band index is not specified, all bands are processed.

Rasters resulting from ST_Clip must have a nodata value assigned for areas clipped, one for each band. If none are provided and the input raster do not have a nodata value defined, nodata values of the resulting raster are set to ST_MinPossibleValue(ST_BandPixelType(rast, band)). When the number of nodata value in the array is smaller than the number of band, the last one in the array is used for the remaining bands. If the number of nodata value is greater than the number of band, the extra nodata values are ignored. All variants accepting an array of nodata values also accept a single value which will be assigned to each band.

If crop is not specified, true is assumed meaning the output raster is cropped to the intersection of the geom and rast extents. If crop is set to false, the new raster gets the same extent as rast.

```

1 -- Clip the first band of an aerial tile by a 20 meter buffer.
2 SELECT ST_Clip(rast, 1,
3   ST_Buffer(ST_Centroid(ST_Envelope(rast)), 20)
4   ) FROM aerials.boston
5 WHERE rid = 4;
6
7 -- Demonstrate effect of crop on final dimensions of raster
8 -- Note how final extent is clipped to that of the geometry
9 -- if crop = true
10 SELECT ST_XMax(ST_Envelope(ST_Clip(rast, 1, clipper, true))) AS xmax_w_trim,
11   ST_XMax(clipper) AS xmax_clipper,
12   ST_XMax(ST_Envelope(ST_Clip(rast, 1, clipper, false))) AS xmax_wo_trim,
13   ST_XMax(ST_Envelope(rast)) AS xmax_rast_orig
14 FROM (SELECT rast, ST_Buffer(ST_Centroid(ST_Envelope(rast)), 6) AS clipper
15      FROM aerials.boston
16      WHERE rid = 6) AS foo;
17
18   xmax_w_trim   |   xmax_clipper   |   xmax_wo_trim   |   xmax_rast_orig
19 -----+-----+-----+-----
20 230657.436173996 | 230657.436173996 | 230666.436173996 | 230666.436173996

```




Full raster tile before clipping



After Clipping

- **ST_DUMPPPOINTS**: A set-returning function (SRF) that extracts the coordinates (vertices) of a geometry. It returns a set of geometry_dump rows, each containing a geometry (geom field) and an array of integers (path field).
 - the geom field POINTs represent the coordinates of the supplied geometry.
 - the path field (an integer[]) is an index enumerating the coordinate positions in the elements of the supplied geometry. The indices are 1-based. For example, for a LINESTRING the paths are {i} where i is the nth coordinate in the LINESTRING. For a POLYGON the paths are {i,j} where i is the ring number (1 is outer; inner rings follow) and j is the coordinate position in the ring.

```

1 SELECT edge_id, (dp).path[1] As index, ST_AsText((dp).geom) As wktnode
2 FROM (SELECT 1 As edge_id
3       , ST_DumpPoints(ST_GeomFromText('LINESTRING(1 2, 3 4, 10 10)')) AS dp
4       UNION ALL
5       SELECT 2 As edge_id
6       , ST_DumpPoints(ST_GeomFromText('LINESTRING(3 5, 5 6, 9 10)')) AS dp
7       ) As foo;
8 edge_id | index |      wktnode
9 -----+-----+-----
10        1 |      1 | POINT(1 2)
11        1 |      2 | POINT(3 4)
12        1 |      3 | POINT(10 10)
13        2 |      1 | POINT(3 5)
14        2 |      2 | POINT(5 6)
15        2 |      3 | POINT(9 10)

```

- **ST_INTERSECTS**: Compares two geometries and returns true if they intersect. Geometries intersect if they have any point in common.

```

1 SELECT ST_Intersects('POINT(0 0)::geometry', 'LINESTRING ( 2 0, 0 2 )::geometry');
2 st_intersects
3 -----
4 f
5 (1 row)
6 SELECT ST_Intersects('POINT(0 0)::geometry', 'LINESTRING ( 0 0, 0 2 )::geometry');
7 st_intersects
8 -----
9 t
10 (1 row)
11
12 -- Look up in table. Make sure table has a GiST index on geometry column for faster
13    lookup.
14 SELECT id, name FROM cities WHERE ST_Intersects(geom, 'SRID=4326;POLYGON((28 53,27.707
15    52.293,27 52,26.293 52.293,26 53,26.293 53.707,27 54,27.707 53.707,28 53))');
16 id | name
17 ---+---
18 2 | Minsk
19 (1 row)

```

- **ST_MAPALGEBRA**: Returns a one-band raster given one or two input rasters, band indexes and one or more user-specified SQL expressions.
 - With one raster: Creates a new one band raster formed by applying a valid PostgreSQL algebraic operation defined by the expression on the input raster (rast). If nband is not provided, band 1 is

assumed. The new raster will have the same georeference, width, and height as the original raster but will only have one band.

- With two rasters: Creates a new one band raster formed by applying a valid PostgreSQL algebraic operation to the two bands defined by the expression on the two input raster bands (rast1, rast2). If no band1, band2 is specified band 1 is assumed. The resulting raster will be aligned (scale, skew and pixel corners) on the grid defined by the first raster. The resulting raster will have the extent defined by the extenttype parameter.

```

1 WITH foo AS (
2     SELECT 1 AS rid, ST_AddBand(ST_AddBand(ST_AddBand(ST_MakeEmptyRaster(2, 2, 0, 0, 1,
3         -1, 0, 0, 0), 1, '16BUI', 1, 0), 2, '8BUI', 10, 0), 3, '32BUI'::text, 100, 0) AS rast
4     UNION ALL
5     SELECT 2 AS rid, ST_AddBand(ST_AddBand(ST_AddBand(ST_MakeEmptyRaster(2, 2, 0, 1, 1,
6         -1, 0, 0, 0), 1, '16BUI', 2, 0), 2, '8BUI', 20, 0), 3, '32BUI'::text, 300, 0) AS rast
7 )
8 SELECT
9     ST_MapAlgebra(
10         t1.rast, 2,
11         t2.rast, 1,
12         '([rast2] + [rast1.val]) / 2'
13     ) AS rast
14 FROM foo t1
15 CROSS JOIN foo t2
16 WHERE t1.rid = 1
17 AND t2.rid = 2;

```

- **ST_RESAMPLE**: Resample a raster using a specified resampling algorithm, new dimensions (width & height), a grid corner (gridx & gridy) and a set of raster georeferencing attributes (scalex, scaley, skewx & skewy) defined or borrowed from another raster. If using a reference raster, the two rasters must have the same SRID. New pixel values are computed using the NearestNeighbor (English or American spelling), Bilinear, Cubic, CubicSpline or Lanczos resampling algorithm. Default is NearestNeighbor which is the fastest but produce the worst interpolation.

```

1 SELECT
2     ST_Width(orig) AS orig_width,
3     ST_Width(reduce_100) AS new_width
4 FROM (
5     SELECT
6         rast AS orig,
7         ST_Resample(rast, 100, 100) AS reduce_100
8     FROM aerials.boston
9     WHERE ST_Intersects(rast,
10         ST_Transform(
11             ST_MakeEnvelope(-71.128, 42.2392, -71.1277, 42.2397, 4326), 26986)
12         )
13     LIMIT 1
14 ) AS foo;
15
16 orig_width | new_width
17 -----+-----
18      200   |      100

```

- **ST_SEGMENTIZE**: Returns a modified geometry having no segment longer than the given max_segment_length. Distance computation is performed in 2d only. For geometry, length units are in units of spatial reference. For geography, units are in meters.

```

1 SELECT ST_AsText(ST_Segmentize(
2     ST_GeomFromText('MULTILINESTRING((-29 -27,-30 -29.7,-36 -31,-45 -33),(-45 -33,-46 -32))'),
3     5)
4 );
5 st_astext
6 --
7 MULTILINESTRING((-29 -27,-30 -29.7,-34.886615700134 -30.758766735029,-36 -31,

```

```

8 -40.8809353009198 -32.0846522890933,-45 -33),
9 (-45 -33,-46 -32))
10 (1 row)
11
12 SELECT ST_AsText(ST_Segmentize(ST_GeomFromText('POLYGON((-29 28, -30 40, -29 28))'),10));
13 st_astext
14 -----
15 POLYGON((-29 28,-29.8304547985374 37.9654575824488,-30 40,-29.1695452014626
16 30.0345424175512,-29 28))
17 (1 row)

```

- **ST_SUMMARYSTATS**: Returns summarystats consisting of count, sum, mean, stddev, min, max for a given raster band of a raster or raster coverage. If no band is specified nband defaults to 1.

```

1 SELECT rid, band, (stats).*
2 FROM (SELECT rid, band, ST_SummaryStats(rast, band) As stats
3       FROM dummy_rast CROSS JOIN generate_series(1,3) As band
4       WHERE rid=2) As foo;

```

rid	band	count	sum	mean	stddev	min	max
2	1	23	5821	253.086957	1.248061	250	254
2	2	25	3682	147.28	59.862188	78	254
2	3	25	3290	131.6	61.647384	62	254

- **ST_UNION**: Unions the input geometries, merging geometry to produce a result geometry with no overlaps. The output may be an atomic geometry, a MultiGeometry, or a Geometry Collection.

```

1 SELECT id,
2       ST_Union(geom) as singlegeom
3 FROM sometable f
4 GROUP BY id;

```

- **ST_VALUE**: Returns the value of a given band in a given columnx, rowy pixel or at a given geometry point. Band numbers start at 1 and band is assumed to be 1 if not specified.

```

1 -- get raster values at particular postgis geometry points
2 -- the srid of your geometry should be same as for your raster
3 SELECT rid, ST_Value(rast, foo.pt_geom) As b1pval, ST_Value(rast, 2, foo.pt_geom) As
4       b2pval
5 FROM dummy_rast CROSS JOIN (SELECT ST_SetSRID(ST_Point(3427927.77, 5793243.76), 0) As
6       pt_geom) As foo
7 WHERE rid=2;

```

rid	b1pval	b2pval
2	252	79

```

11
12 -- general fictitious example using a real table
13 SELECT rid, ST_Value(rast, 3, sometable.geom) As b3pval
14 FROM sometable
15 WHERE ST_Intersects(rast,sometable.geom);

```